

IMDb RNN Experiments — Summary

Configuration

Data: cleaned, tokenized, top 10k vocabulary from IMDb 50k with a predetermined 25k/25k split; IDs padded to {25, 50, 100}.

Model Template: Embedding (100) -> {RNN/LSTM/BiLSTM, 2 layers, hidden=64, dropout=0.5} -> mean-pool activation -> FC(1) is the model template.

Loss: BCEWithLogits; Device=CPU; Batch=32; Epochs=5; Seed=1337.

Metrics: Macro-F1, accuracy, and training time/epoch.

Results

	model	activation	optimizer	seq_len	grad_clip	epochs	final_test_loss	accuracy	f1_macro	epoch_time_s(last_train)	device	seed
0	RNN	relu	ADAM	50	off	5	0.480314	0.76792	0.767447	1.6807	cpu	1337
1	LSTM	relu	ADAM	50	off	5	0.490834	0.76404	0.764038	6.3161	cpu	1337
2	BILSTM	relu	ADAM	50	off	5	0.487694	0.76188	0.761685	13.6662	cpu	1337
3	LSTM	sigmoid	ADAM	50	off	5	0.508847	0.75428	0.753633	6.6786	cpu	1337
4	LSTM	relu	ADAM	50	off	5	0.481788	0.76440	0.763800	6.5080	cpu	1337
5	LSTM	tanh	ADAM	50	off	5	0.484845	0.76512	0.764925	6.0606	cpu	1337
6	LSTM	relu	ADAM	50	off	5	0.488589	0.75704	0.757032	6.4766	cpu	1337
7	LSTM	relu	SGD	50	off	5	0.692958	0.49916	0.334303	6.0823	cpu	1337
8	LSTM	relu	RMSPROP	50	off	5	0.486625	0.75968	0.757985	6.1206	cpu	1337
9	LSTM	relu	ADAM	25	off	5	0.559067	0.70484	0.704575	3.1427	cpu	1337
10	LSTM	relu	ADAM	50	off	5	0.484411	0.76384	0.763437	6.2108	cpu	1337
11	LSTM	relu	ADAM	100	off	5	0.396150	0.82272	0.822396	12.3214	cpu	1337
12	LSTM	relu	ADAM	50	off	5	0.501528	0.75160	0.750061	6.2651	cpu	1337
13	LSTM	relu	ADAM	50	1.0	5	0.493700	0.76196	0.760032	6.3978	cpu	1337

What worked the best

Overall best: LSTM, ReLU, Adam, seq_len=100 => Acc 0.823, F1= 0.822, approx. 12.32 s/epoch.
In five epochs, this configuration is the only one that pushes >0.82 on CPU.

Factor-by-factor effects (changing one variable at a time)

Architecture (no clipping, Adam, ReLU, seq_len=50)

RNN: Acc=0.768 / F1=0.767 / approx. 1.68 s/epoch (fastest).

LSTM: approx. 6.32 s, Acc 0.764, and F1=0.764.

BiLSTM: Acc=0.762, F1=0.762, approx. 13.67 s (almost 2 x slower; no gain).

Conclusion: LSTM is equivalent to RNN on short sequences (50), however BiLSTM's additional expense was ineffective.

Activation (len=50, LSTM, Adam)

ReLU: Acc=0.764/F1=0.764 (baseline).

Sigmoid: F1=0.754 (lower) and Acc=0.754.

Tanh: Acc=0.765/F1=0.765 (equivalent to ReLU).

Conclusion: Sigmoid performs poorly as a hidden activation here, although ReLU/tanh are safe.

Optimizer (len=50, LSTM, ReLU)

Adam: F1=0.764/ Acc=0.764 (steady).

RMSProp: somewhat worse Acc=0.760/F1=0.758.

SGD (default LR): Acc=0.499 / F1=0.334 (collapsed).

Conclusion: Adam is the ideal plug-and-play option. The SGD run probably utilized an LR that was too small for this configuration; it would be competitive with a tweaked LR (e.g., 0.01).

Length of sequence (LSTM, ReLU, Adam)

len=25: Acc=0.705/F1=0.705/3.14 s.

len=50: Acc=0.764 / F1=0.763 / 6.21 s.

len=100: F1=0.822 / Acc=0.823 / 12.32 s.

Conclusion: Cost scales roughly linearly with length; accuracy/F1 clearly improves with greater context.

Gradient clipping (len=50, LSTM, ReLU, Adam)

Off: F= 0.763, Acc= 0.764.

max_norm=1.0: F= 0.760, Acc=0.762.

Conclusion: In this stable regime (5 epochs, CPU), clipping did not significantly alter the results.

Observations during training

At len=50, BiLSTM $\approx 2 \times$ LSTM time.

With significant accuracy gains, doubling seq_len (50->100) almost doubles epoch time (6.2->12.3 s).

At larger lengths, accuracy lags behind LSTM, although plain RNN is the fastest.

Recommendations

- i) LSTM + ReLU + Adam is the default for reports and downstream tasks.
- ii) Use seq_len=100 and keep epochs small (≥ 5) if you want the highest CPU accuracy.
- iii) Seq_len=50 is a nice balance (approx. 6 s/epoch) with approx. 0.764 F1 if you need quick iterations.
- iv) Here, stay away from BiLSTM unless you can take use of the additional expense (longer training, GPU, etc.).
- v) ReLU or tanh are preferable to sigmoid hidden activations.
- vi) LR tweaking is necessary for SGD; it collapsed with default settings. Rerun SGD with lr equivalent to 1e-2 and momentum 0.9 if you need to compare fairly.